

# Machine Learning and Data Mining

Introduction

Albert Bifet(@abifet)



# Data Science

**Data Science** is an interdisciplinary field focused on extracting knowledge or insights from large volumes of data.

# Data Scientist



Figure: <http://www.marketingdistillery.com/2014/11/29/is-data-science-a-buzzword-modern-data-scientist-defined/>

# Data Science




Figure: Drew Conway's Venn diagram

# Classification

## Definition

Given  $n_C$  different classes, a classifier algorithm builds a model that predicts for every unlabelled instance  $I$  the class  $C$  to which it belongs with accuracy.

## Example

A spam filter

## Example

Twitter Sentiment analysis: analyze tweets with positive or negative feelings

# Classification

Data set that describes e-mail features for deciding if it is spam.

## Example

<b>Contains "Money"</b>	<b>Domain type</b>	<b>Has attach.</b>	<b>Time received</b>	<b>spam</b>
yes	com	yes	night	yes
yes	edu	no	night	yes
no	com	yes	night	yes
no	edu	no	day	no
no	com	no	day	no
yes	cat	no	day	yes

Assume we have to classify the following new instance:

<b>Contains "Money"</b>	<b>Domain type</b>	<b>Has attach.</b>	<b>Time received</b>	<b>spam</b>
yes	edu	yes	day	?

# $k$ -Nearest Neighbours

## $k$ -NN Classifier

- Training: store all instances in memory
- Prediction:
  - Find the  $k$  nearest instances
  - Output majority class of these  $k$  instances

# Bayes Classifiers

## Naïve Bayes

- Based on Bayes Theorem:

$$P(c|d) = \frac{P(c)P(d|c)}{P(d)}$$

$$\text{posterior} = \frac{\text{prior} \times \text{likelihood}}{\text{evidence}}$$

- Estimates the probability of observing attribute  $a$  and the prior probability  $P(c)$
- Probability of class  $c$  given an instance  $d$ :

$$P(c|d) = \frac{P(c) \prod_{a \in d} P(a|c)}{P(d)}$$




# Bayes Classifiers

## Multinomial Naïve Bayes

- Considers a document as a bag-of-words.
- Estimates the probability of observing word  $w$  and the prior probability  $P(c)$
- Probability of class  $c$  given a test document  $d$ :


$$P(c|d) = \frac{P(c) \prod_{w \in d} P(w|c)^{n_{wd}}}{P(d)}$$

# Perceptron



- Data stream:  $\langle \vec{x}_i, y_i \rangle$
- Classical perceptron:  $h_{\vec{w}}(\vec{x}_i) = \text{sgn}(\vec{w}^T \vec{x}_i)$ ,
- Minimize Mean-square error:  $J(\vec{w}) = \frac{1}{2} \sum (y_i - h_{\vec{w}}(\vec{x}_i))^2$

# Perceptron



- We use sigmoid function  $h_{\vec{w}} = \sigma(\vec{w}^T \vec{x})$  where

$$\sigma(x) = 1/(1 + e^{-x})$$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

# Perceptron

- Minimize Mean-square error:  $J(\vec{w}) = \frac{1}{2} \sum (y_i - h_{\vec{w}}(\vec{x}_i))^2$
- Stochastic Gradient Descent:  $\vec{w} = \vec{w} - \eta \nabla J \vec{x}_i$
- Gradient of the error function:


$$\nabla J = - \sum_i (y_i - h_{\vec{w}}(\vec{x}_i)) \nabla h_{\vec{w}}(\vec{x}_i)$$

$$\nabla h_{\vec{w}}(\vec{x}_i) = h_{\vec{w}}(\vec{x}_i)(1 - h_{\vec{w}}(\vec{x}_i))$$

- Weight update rule

$$\vec{w} = \vec{w} + \eta \sum_i (y_i - h_{\vec{w}}(\vec{x}_i)) h_{\vec{w}}(\vec{x}_i)(1 - h_{\vec{w}}(\vec{x}_i)) \vec{x}_i$$

# Restricted Boltzmann Machines (RBMs)



- Energy-based models, where

$$P(\vec{x}, \vec{z}) \propto e^{-E(\vec{x}, \vec{z})}.$$

- Manipulate a weight matrix  $W$  to find low-energy states and thus generate high probability  $P(\vec{x}, \vec{z})$ , where

$$E(\vec{x}, \vec{z}) = -W.$$

- RBMs can be stacked on top of each other to form so-called **Deep Belief Networks (DBNs)**

# Classification

Data set that describes e-mail features for deciding if it is spam.

## Example

<b>Contains "Money"</b>	<b>Domain type</b>	<b>Has attach.</b>	<b>Time received</b>	<b>spam</b>
yes	com	yes	night	yes
yes	edu	no	night	yes
no	com	yes	night	yes
no	edu	no	day	no
no	com	no	day	no
yes	cat	no	day	yes


Assume we have to classify the following new instance:

<b>Contains "Money"</b>	<b>Domain type</b>	<b>Has attach.</b>	<b>Time received</b>	<b>spam</b>
yes	edu	yes	day	?

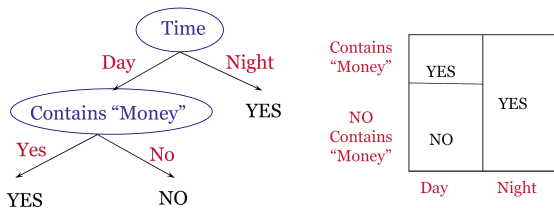
# Classification

- Assume we have to classify the following new instance:

<b>Contains "Money"</b>	<b>Domain type</b>	<b>Has attach.</b>	<b>Time received</b>	<b>spam</b>
yes	edu	yes	day	?



# Decision Trees



Basic induction strategy:

- $A \leftarrow$  the "best" decision attribute for next *node*
- Assign  $A$  as decision attribute for *node*
- For each value of  $A$ , create new descendant of *node*
- Sort training examples to leaf nodes
- If training examples perfectly classified, Then STOP, Else iterate over new leaf nodes



# Bagging

## Example

Dataset of 4 Instances : A, B, C, D

Classifier 1: B, A, C, B

Classifier 2: D, B, A, D

Classifier 3: B, A, C, B

Classifier 4: B, C, B, B

Classifier 5: D, C, A, C

Bagging builds a set of  $M$  base models, with a bootstrap sample created by drawing random samples with replacement.

# Random Forests

- Bagging
- Random Trees: trees that in each node only uses a random subset of the attributes

Random Forests is one of the most popular methods in machine learning.

# Boosting

## The strength of Weak Learnability, Schapire 90

A boosting algorithm transforms a weak learner  
into a strong one

# Boosting

## A formal description of Boosting (Schapire)

- given a training set  $(x_1, y_1), \dots, (x_m, y_m)$
- $y_i \in \{-1, +1\}$  correct label of instance  $x_i \in X$
- for  $t = 1, \dots, T$ 
  - construct distribution  $D_t$
  - find weak classifier

$$h_t : X \rightarrow \{-1, +1\}$$

with small error  $\varepsilon_t = \Pr_{D_t}[h_t(x_i) \neq y_i]$  on  $D_t$

- output final classifier

# Boosting

## AdaBoost

- 1: Initialize  $D_1(i) = 1/m$  for all  $i \in \{1, 2, \dots, m\}$
- 2: **for**  $t = 1, 2, \dots, T$  **do**
- 3:   Call **WeakLearn**, providing it with distribution  $D_t$
- 4:   Get back hypothesis  $h_t : X \rightarrow Y$
- 5:   Calculate error of  $h_t$  :  $\varepsilon_t = \sum_{i: h_t(x_i) \neq y_i} D_t(i)$
- 6:   Update distribution

$$D_t : D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} \varepsilon_t / (1 - \varepsilon_t) & \text{if } h_t(x_i) = y_i \\ 1 & \text{otherwise} \end{cases}$$

where  $Z_t$  is a normalization constant (chosen so  $D_{t+1}$  is a probability distribution)

- 7: **return**  $h_{fin}(x) = \arg \max_{y \in Y} \sum_{t: h_t(x) = y} -\log \varepsilon_t / (1 - \varepsilon_t)$

# Boosting

## AdaBoost

- 1: Initialize  $D_1(i) = 1/m$  for all  $i \in \{1, 2, \dots, m\}$
- 2: **for**  $t = 1, 2, \dots, T$  **do**
- 3:   Call **WeakLearn**, providing it with distribution  $D_t$
- 4:   Get back hypothesis  $h_t : X \rightarrow Y$
- 5:   Calculate error of  $h_t$  :  $\varepsilon_t = \sum_{i:h_t(x_i) \neq y_i} D_t(i)$
- 6:   Update distribution

$$D_t : D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} \varepsilon_t & \text{if } h_t(x_i) = y_i \\ 1 - \varepsilon_t & \text{otherwise} \end{cases}$$

where  $Z_t$  is a normalization constant (chosen so  $D_{t+1}$  is a probability distribution)

- 7: **return**  $h_{fin}(x) = \arg \max_{y \in Y} \sum_{t:h_t(x)=y} -\log \varepsilon_t / (1 - \varepsilon_t)$

# Stacking

Use a classifier to combine predictions of base classifiers

## Example

- Use a perceptron to do stacking
- Use decision trees as base classifiers

# Clustering

## Definition

Clustering is the distribution of a set of instances of examples into non-known groups according to some common relations or affinities.

## Example

Market segmentation of customers

## Example

Social network communities



# Clustering

## Definition

Given

- a set of instances  $I$
- a number of clusters  $K$
- an objective function  $\text{cost}(I)$

a clustering algorithm computes an assignment of a cluster for each instance

$$f : I \rightarrow \{1, \dots, K\}$$

that minimizes the objective function  $\text{cost}(I)$

# Clustering

## Definition

Given

- a set of instances  $I$
- a number of clusters  $K$
- an objective function  $\text{cost}(C, I)$

a clustering algorithm computes a set  $C$  of instances with  $|C| = K$  that minimizes the objective function

$$\text{cost}(C, I) = \sum_{x \in I} d^2(x, C)$$

where

- $d(x, c)$ : distance function between  $x$  and  $c$
- $d^2(x, C) = \min_{c \in C} d^2(x, c)$ : distance from  $x$  to the nearest point in  $C$

# k-means

- 1. Choose  $k$  initial centers  $C = \{c_1, \dots, c_k\}$
- 2. while stopping criterion has not been met
  - For  $i = 1, \dots, N$ 
    - find closest center  $c_k \in C$  to each instance  $p_i$
    - assign instance  $p_i$  to cluster  $C_k$
  - For  $k = 1, \dots, K$ 
    - set  $c_k$  to be the center of mass of all points in  $C_i$

# k-means++

- 1. Choose a initial center  $c_1$
- For  $k = 2, \dots, K$ 
  - select  $c_k = p \in I$  with probability  $d^2(p, C)/\text{cost}(C, I)$
- 2. while stopping criterion has not been met
  - For  $i = 1, \dots, N$ 
    - find closest center  $c_k \in C$  to each instance  $p_i$
    - assign instance  $p_i$  to cluster  $C_k$
  - For  $k = 1, \dots, K$ 
    - set  $c_k$  to be the center of mass of all points in  $C_i$

# Performance Measures

## Internal Measures

- Sum square distance
- Dunn index  $D = \frac{d_{min}}{d_{max}}$
- C-Index  $C = \frac{S - S_{min}}{S_{max} - S_{min}}$

## External Measures

- Rand Measure
- F Measure
- Jaccard
- Purity

# Density based methods

## DBSCAN

- $\varepsilon$ -neighborhood( $p$ ): set of points that are at a distance of  $p$  less or equal to  $\varepsilon$
- Core object: object whose  $\varepsilon$ -neighborhood has an overall weight at least  $\mu$
- A point  $p$  is *directly density-reachable* from  $q$  if
  - $p$  is in  $\varepsilon$ -neighborhood( $q$ )
  - $q$  is a core object
- A point  $p$  is *density-reachable* from  $q$  if
  - there is a chain of points  $p_1, \dots, p_n$  such that  $p_{i+1}$  is directly density-reachable from  $p_i$
- A point  $p$  is *density-connected* from  $q$  if
  - there is point  $o$  such that  $p$  and  $q$  are density-reachable from  $o$

# Density based methods

## DBSCAN

- A *cluster*  $C$  of points satisfies
  - if  $p \in C$  and  $q$  is density-reachable from  $p$ , then  $q \in C$
  - all points  $p, q \in C$  are density-connected
- A *cluster* is uniquely determined by any of its core points
- A *cluster* can be obtained
  - choosing an arbitrary core point as a seed
  - retrieve all points that are density-reachable from the seed

# DBSCAN

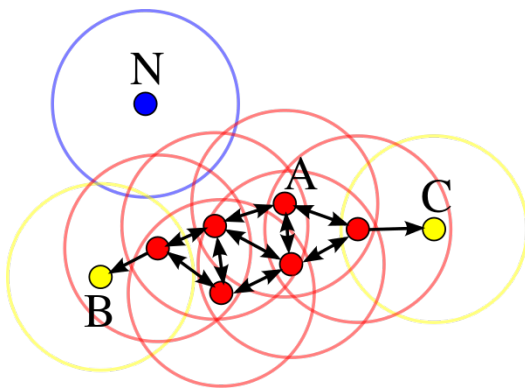


Figure: DBSCAN Point Example with  $\mu=3$



# Density based methods

## DBSCAN

- select an arbitrary point  $p$
- retrieve all points density-reachable from  $p$
- if  $p$  is a core point, a cluster is formed
- If  $p$  is a border point
  - no points are density-reachable from  $p$
  - DBSCAN visits the next point of the database
- Continue the process until all of the points have been processed

# Frequent Patterns

Suppose  $\mathcal{D}$  is a dataset of patterns,  $t \in \mathcal{D}$ , and  $min\_sup$  is a constant.

## Definition

*Support* ( $t$ ): number of patterns in  $\mathcal{D}$  that are superpatterns of  $t$ .

## Definition

Pattern  $t$  is *frequent* if  $Support(t) \geq min\_sup$ .

## Frequent Subpattern Problem

Given  $\mathcal{D}$  and  $min\_sup$ , find all frequent subpatterns of patterns in  $\mathcal{D}$ .

# Frequent Patterns

Suppose  $\mathcal{D}$  is a dataset of patterns,  $t \in \mathcal{D}$ , and  $min\_sup$  is a constant.

## Definition

**Support** ( $t$ ): number of patterns in  $\mathcal{D}$  that are superpatterns of  $t$ .

## Definition

Pattern  $t$  is **frequent** if  $Support(t) \geq min\_sup$ .

## Frequent Subpattern Problem

Given  $\mathcal{D}$  and  $min\_sup$ , find all frequent subpatterns of patterns in  $\mathcal{D}$ .

# Frequent Patterns

Suppose  $\mathcal{D}$  is a dataset of patterns,  $t \in \mathcal{D}$ , and  $min\_sup$  is a constant.

## Definition

**Support** ( $t$ ): number of patterns in  $\mathcal{D}$  that are superpatterns of  $t$ .

## Definition

Pattern  $t$  is **frequent** if  $Support(t) \geq min\_sup$ .

## Frequent Subpattern Problem

Given  $\mathcal{D}$  and  $min\_sup$ , find all frequent subpatterns of patterns in  $\mathcal{D}$ .

# Frequent Patterns

Suppose  $\mathcal{D}$  is a dataset of patterns,  $t \in \mathcal{D}$ , and  $min\_sup$  is a constant.

## Definition

**Support** ( $t$ ): number of patterns in  $\mathcal{D}$  that are superpatterns of  $t$ .

## Definition

Pattern  $t$  is **frequent** if  $Support(t) \geq min\_sup$ .

## Frequent Subpattern Problem

Given  $\mathcal{D}$  and  $min\_sup$ , find all frequent subpatterns of patterns in  $\mathcal{D}$ .

# Pattern Mining

## Dataset Example

Document	Patterns
d1	abce
d2	cde
d3	abce
d4	acde
d5	abcde
d6	bcd

# Itemset Mining

d1	abce
d2	cde
d3	abce
d4	acde
d5	abcde
d6	bcd

Support	Frequent
d1,d2,d3,d4,d5,d6	c
d1,d2,d3,d4,d5	e,ce
d1,d3,d4,d5	a,ac,ae,ace
d1,d3,d5,d6	b,bc
d2,d4,d5,d6	d,cd
d1,d3,d5	ab,abc,abe be,bce,abce
d2,d4,d5	de,cde

minimal support = 3

# Itemset Mining

d1 abce  
d2 cde  
d3 abce  
d4 acde  
d5 abcde  
d6 bcd

Support	Frequent
6	c
5	e,ce
4	a,ac,ae,ace
4	b,bc
4	d,cd
3	ab,abc,abe be,bce,abce
3	de,cde



# Itemset Mining

d1 abce  
d2 cde  
d3 abce  
d4 acde  
d5 abcde  
d6 bcd

Support	Frequent	Gen	Closed
6	c	c	c
5	e,ce	e	ce
4	a,ac,ae,ace	a	ace
4	b,bc	b	bc
4	d,cd	d	cd
3	ab,abc,abe be,bce,abce	ab be	abce
3	de,cde	de	cde

# Itemset Mining

		<b>Support</b>	<b>Frequent</b>	<b>Gen</b>	<b>Closed</b>	<b>Max</b>
d1	abce	6	c	c	c	
d2	cde	5	e,ce	e	ce	
d3	abce	4	a,ac,ae,ace	a	ace	
d4	acde	4	b,bc	b	bc	
d5	abcde	4	d,cd	d	cd	
d6	bcd	3	ab,abc,abe	ab		
			be,bce,abce	be	abce	abce
		3	de,cde	de	cde	cde

# Itemset Mining

		Support	Frequent	Gen	Closed	Max
d1	ab <b>c</b> e	6	<b>c</b>	<b>c</b>	<b>c</b>	
d2	<b>c</b> de	5	e,ce	e	ce	
d3	ab <b>c</b> e	4	a,ac,ae,ace	a	ace	
d4	a <b>c</b> de	4	b,bc	b	bc	
d5	ab <b>c</b> de	4	d,cd	d	cd	
d6	b <b>c</b> d	3	ab,abc,abe be,bce,abce	ab be	abce	abce
		3	de,cde	de	cde	cde

# Itemset Mining

d1 ab**c**de  
 d2 **c**de  
 d3 ab**c**de  
 d4 a**c**de  
 d5 ab**c**de  
 d6 bcd

e → ce

Support	Frequent	Gen	Closed	Max
6	c	c	c	
5	e,ce	e	ce	
4	a,ac,ae,ace	a	ace	
4	b,bc	b	bc	
4	d,cd	d	cd	
3	ab,abc,abe	ab		
	be,bce,abce	be	abce	abce
3	de,cde	de	cde	cde

# Itemset Mining

		<b>Support</b>	<b>Frequent</b>	<b>Gen</b>	<b>Closed</b>	<b>Max</b>
d1	ab <b>ce</b>	6	c	c	c	
d2	<b>cde</b>	5	e,ce	e	ce	
d3	ab <b>ce</b>	4	a,ac,ae,ace	a	ace	
d4	a <b>cde</b>	4	b,bc	b	bc	
d5	ab <b>cde</b>	4	d,cd	d	cd	
d6	bcd	3	ab,abc,abe be,bce,abce	ab be	abce	abce
		3	de,cde	de	cde	cde

# Itemset Mining

		<b>Support</b>	<b>Frequent</b>	<b>Gen</b>	<b>Closed</b>	<b>Max</b>
d1	abce	6	c	c	c	
d2	cde	5	e,ce	e	ce	
d3	abce	4	a,ac,ae,ace	a	ace	
d4	acde	4	b,bc	b	bc	
d5	abcde	4	d,cd	d	cd	
d6	bcd	3	ab,abc,abe	ab		
			be,bce,abce	be	abce	abce
		3	de,cde	de	cde	cde

# Itemset Mining

		Support	Frequent	Gen	Closed	Max
d1	abce	6	c	c	c	
d2	cde	5	e,ce	e	ce	
d3	abce	4	a,ac,ae,ace	a	ace	
d4	acde	4	b,bc	b	bc	
d5	abcde	4	d,cd	d	cd	
d6	bcd	3	ab,abc,abe	ab		
	a → ace		be,bce,abce	be	abce	abce
		3	de,cde	de	cde	cde

# Itemset Mining

		<b>Support</b>	<b>Frequent</b>	<b>Gen</b>	<b>Closed</b>	<b>Max</b>
d1	abce	6	c	c	c	
d2	cde	5	e,ce	e	ce	
d3	abce	4	a,ac,ae,ace	a	ace	
d4	acde	4	b,bc	b	bc	
d5	abcde	4	d,cd	d	cd	
d6	bcd	3	ab,abc,abe be,bce,abce	ab be	abce	abce
		3	de,cde	de	cde	cde



# Closed Patterns

Usually, there are too many frequent patterns. We can compute a smaller set, while keeping the same information.

## Example

A set of 1000 items, has  $2^{1000} \approx 10^{301}$  subsets, that is more than the number of atoms in the universe  $\approx 10^{79}$

# Closed Patterns

## *A priori* property

If  $t'$  is a subpattern of  $t$ , then  $\text{Support}(t') \geq \text{Support}(t)$ .

## Definition

A frequent pattern  $t$  is *closed* if none of its proper superpatterns has the same support as it has.

Frequent subpatterns and their supports can be generated from closed patterns.

# Maximal Patterns

## Definition

A frequent pattern  $t$  is *maximal* if none of its proper superpatterns is frequent.

Frequent subpatterns can be generated from maximal patterns, but not with their support.

All maximal patterns are closed, but not all closed patterns are maximal.

# Non streaming frequent itemset miners

## Representation:

- Horizontal layout

T1: a, b, c

T2: b, c, e

T3: b, d, e

- Vertical layout

a: 1 0 0

b: 1 1 1

c: 1 1 0

## Search:

- Breadth-first (levelwise): Apriori
- Depth-first: Eclat, FP-Growth

# The Apriori Algorithm

## APRIORI ALGORITHM

- 1 Initialize the item set size  $k = 1$
- 2 Start with single element sets
- 3 Prune the non-frequent ones
- 4 **while** there are frequent item sets
- 5     **do** create candidates with one item more
- 6         Prune the non-frequent ones
- 7         Increment the item set size  $k = k + 1$
- 8 Output: the frequent item sets

# The Eclat Algorithm

## Depth-First Search

- divide-and-conquer scheme : the problem is processed by splitting it into smaller subproblems, which are then processed recursively
  - **conditional database for the prefix a**
    - transactions that contain a
  - **conditional database for item sets without a**
    - transactions that not contain a
- Vertical representation
- Support counting is done by intersecting lists of transaction identifiers

# The FP-Growth Algorithm

## Depth-First Search

- divide-and-conquer scheme : the problem is processed by splitting it into smaller subproblems, which are then processed recursively
  - **conditional database for the prefix a**
    - transactions that contain a
  - **conditional database for item sets without a**
    - transactions that not contain a
- Vertical and Horizontal representation : FP-Tree
  - prefix tree with links between nodes that correspond to the same item
- Support counting is done using FP-Tree

# Mining Graph Data

## Problem

Given a data set  $\mathcal{D}$  of graphs, find frequent graphs.

Transaction Id	Graph
1	$\begin{array}{c} \text{O} \\   \\ \text{C} - \text{C} - \text{S} - \text{N} \\   \\ \text{O} \end{array}$
2	$\begin{array}{c} \text{O} \\   \\ \text{C} - \text{C} - \text{S} - \text{N} \\   \\ \text{C} \end{array}$
3	$\begin{array}{c} \text{N} \\    \\ \text{C} - \text{C} - \text{S} - \text{N} \end{array}$



# The gSpan Algorithm

$\text{GSPAN}(g, D, \text{min\_sup}, S)$

Input: A graph  $g$ , a graph dataset  $D$ ,  $\text{min\_sup}$ .

Output: The frequent graph set  $S$ .

```
1  if  $g \neq \text{min}(g)$ 
2    then return  $S$ 
3  insert  $g$  into  $S$ 
4  update support counter structure
5   $C \leftarrow \emptyset$ 
6  for each  $g'$  that can be right-most
    extended from  $g$  in one step
7    do if  $\text{support}(g) \geq \text{min\_sup}$ 
8      then insert  $g'$  into  $C$ 
9  for each  $g'$  in  $C$ 
10    do  $S \leftarrow \text{GSPAN}(g', D, \text{min\_sup}, S)$ 
11  return  $S$ 
```